

EVENT DETECTION AND CLASSIFICATION

Sameer Chauhan, Sharang Phadke, Christian Sherland

Cooper Union for the Advancement of Science and Art
Electrical Engineering
41 Cooper Square
New York, NY 10003

ABSTRACT

The IEEE AASP Challenge addresses the problem of acoustic event detection and classification in an office environment. Our system performs segmentation and event classification on a continuous stream of acoustic activity in an office using basic feature extraction techniques and a two layer classifier (we haven't implemented two layers yet).

Index Terms—

1. INTRODUCTION

The task of event detection and classification is fundamentally important in computational auditory scene analysis (CASA). Detecting events such as words in a sentence of speech, the arrival of a bus or train, or any other common occurrence in an audio stream is the first step of any speech or audio processing system that needs to make use of this information. Many such systems then need to classify the detected events in order to take appropriate actions.

In this project submission, the problem of event detection and classification, as applied to an office environment, was approached from a pattern recognition perspective. Our system consists of a segmentation stage, a feature extraction stage, and two classification stages. The first stage detects the onset and offset of events in a live recording, the second extracts features from each event that can be used to classify events, and the final stages classify each detected event using two pre-trained classifiers. The classification system initially classifies the detected event as one of a few groups of events. The system then classifies the events in each group as a specific type of event using a set of features that extracts the most discriminatory information from the events within each group.

2. FEATURE EXTRACTION

In order to classify events, our system extracts a set of features from the training data. This set includes the spectral centroid, spectral flux, spectral sparsity, temporal sparsity, loudness, short time energy and Mel-frequency cepstrum coefficients (MFCC). The MFCCs are based on frequency bands equally spaced on the mel scale which

approximates the human auditory system's response. Each feature highlights different acoustic properties of the signal. Spectral sparsity is expected to be very large for pure sine tones or bells and smaller for sounds with significant "noise" characteristics that imply a wide frequency spectrum. Temporal sparsity is large for sounds such as footsteps in relative silence and is useful for indexing and retrieving these types of sounds. Short time energy is analogous to the volume of the event and the spectral centroid is essentially the center of mass of the spectrum. They are both expected to be reliable indicators of silence. The spectral flux, also called spectral variation, measures how quickly the power spectrum changes. It can be used to determine the timbre of the audio signal. Different permutations of these features will be tested and the best combination will be used for the system training and classification stages of our system.

Our feature set is computed within a 40ms window with no overlap, which forces our frames to be 40ms long. This limits our precision but increases our accuracy. In order to increase the performance of the system, we will choose features to be calculated over a long period (1s). This 1s window will slide every 20ms and overlap will be included with the features obtained within the 40ms window to help obtain a more precise feature set for the events.

3. TRAINING

The classification stage of our system employs two levels of classification. As input the classifier takes in an event signal and first classifies what type of event it is. Then based upon the type of event the signal is passed through a type-specific classifier which determines the exact event contained in the signal. In the second stage each type has its own classifier.

Each level of classification in the system was trained with features extracted from the provided training data. The first stage classifier was trained using the full set of extracted features for all events. The second stage classifiers were each trained using features extracted only from training signals of the correct type. Additionally, the features used for training were chosen to best describe the type of event being classified.

4. SEGMENTATION

Before events within an input signal can be classified, the segments of the input which contain events must be extracted. This is done by examining the short time energy and spectral centroid features. The segmenter examines peaks of the features and checks them against a threshold to determine the onset and offset of events in the signal. The onset and offset times are then used to extract the portions of

Thanks to the Cooper Union

This document is licensed under the Creative Commons Attribution 3.0 License (CC BY 3.0).

<http://creativecommons.org/licenses/by/3.0/>

© 2013 The Authors.

the input which contain an event. Extracted portions are then passed to the classification stage.

5. CLASSIFICATION

After a set of event signals has been segmented, the classifier extracts a full set of features for each event signal. For each event, these features are passed into the first level classifier. Based upon the output of the first stage classifier, each event is labeled with its subgroup.

Each event signal was then passed through the classifier that corresponds to its subgroup label and reclassified based upon the subset of features that best describes its type of event.

Types were chosen such that members of a given type generally have many similar features. The set of features that is used to determine the specific event that occurred in stage two is the set which varies most among the members.

6. RESULTS

In order to evaluate the performance of our system we examined the percentage of frames correctly classified with our segmentation and with perfect segmentation. The perfect segmentation and ground truth label for each frame were determined using the annotation to the development data.

The percentage of frames correct given perfect segmentation gives a sense of how well the classification stage of our system performs. Currently, given perfect segmentation, our system correctly classifies 48% of frames.

The percentage of frames correct with segmentation determined by our system when compared to the other metric gives a good sense of how segmentation effects classification. With our segmentation, the system correctly classified 25% of frames.

The confusion matrices of classifier stages of the perfectly segmented events, as well as the events segmented using our algorithm can be seen in Figure 1 and Figure 2 respectively. (We need to make the figure bigger...)

		Classifier Confusion Matrix Without Segmenter																	
Truth	alert	0	0	4.9	28	9.8	9.8	0	12.2	0	6.1	12.2	17.1	0	0				[82]
	clearthroat	0	70.8	4.2	0	20.8	0	0	4.2	0	0	0	0	0	0				[24]
	cough	0	2	58.2	1	23.5	0	0	13.3	0	0	0	2	0	0				[98]
	knock	0	0	0	60.7	8.9	0	0	0	0	0	0	28.6	1.8	0				[56]
	drawer	0	1.2	1.2	0	69.3	0	0	0	0	0	0	8.3	0	0				[84]
	keyboard	0	0	0	0	18.8	43.8	1	0	6.3	1	0	27.1	0	0				[96]
	keys	0	0	0.7	0	9.2	6.3	51.4	1.4	1.4	3.5	0	25.4	0	0.7				[142]
	laughter	0	0	12.5	8.3	12.5	8.3	0	50	0	0	0	8.3	0	0				[24]
	pageturn	0	0	0	0	24.7	6.2	37.1	1	6.2	15.5	0	9.3	0	0				[97]
	pendrop	0	0	2.7	0	48.6	0	2.7	0	0	10.8	0	35.1	0	0				[37]
	phone	0	0	4.5	0	1.5	0	0	6	0	0	85.1	3	0	0				[67]
	printer	0	0	0	0	0.8	0.8	0	0	0	4.2	0	94.2	0	0				[120]
	speech	0	10.1	11.2	6.7	13.5	0	6.7	33.7	1.1	1.1	0	10.1	5.6	0				[89]
	switch	0	0	0	0	0	0	0	28.6	0	0	0	57.1	0	14.3				[7]
		alert	clearthroat	cough	knock	drawer	keyboard	keys	laughter	pageturn	pendrop	phone	printer	speech	switch				
		Response																	

Figure 1: Example of a figure with experimental results.

		Classifier Confusion Matrix With Segmenter																	
Truth	unknown	0	0	0	0	21.4	0	0	1.4	0	0	77.1	0						[70]
	alert	0	0	0	7.3	22	0	4.9	2.4	2.4	14.6	46.3	0						[41]
	clearthroat	0	0	26.7	0	61.7	0	1.7	0	0	0	10	0						[60]
	cough	0	0	0.4	22.4	61.8	0.4	7.1	0	0	0	7.9	0						[241]
	drawer	0	0	1.3	1.3	76	1.3	1.3	0	0	1.3	17.3	0						[75]
	keys	0	0	0	0	26.5	56.6	0.9	1.8	0.9	0	13.3	0						[113]
	laughter	0	0	0	8.3	47.2	0	33.3	0	0	0	11.1	0						[36]
	pageturn	0	0	0	0	12.9	37.1	0	11.3	14.5	0	24.2	0						[62]
	pendrop	0	0	0	0	58.9	3.2	0	3.2	4.2	0	30.5	0						[95]
	phone	0	0	0	4.3	2.9	0	5.7	0	0	80	7.1	0						[70]
	printer	0	0	0	0	0	0	0	31.2	0	0	68.8	0						[16]
	speech	0	0	11.3	11.3	18.3	1.4	36.6	0	1.4	0	15.5	4.2						[71]
		unknown	alert	clearthroat	cough	drawer	keys	laughter	pageturn	pendrop	phone	printer	speech						
		Response																	

Figure 2: Example of a figure with experimental results.

7. CONCLUSION

8. REFERENCES

List and number all bibliographical references at the end of the paper. The references should be numbered in order of appearance in the document. When referring to them in the text, type the corresponding reference number in square brackets as shown at the end of this sentence [?], [?]. For \LaTeX users, the use of the Bib \TeX style file IEEEtran.bst is recommended, which is included in the \LaTeX paper kit available from the workshop website [?].

9. ACKNOWLEDGMENT

The preferred spelling of the word acknowledgment in America is without an “e” after the “g.” Try to avoid the stilted expression, “One of us (R. B. G.) thanks ...” Instead, try “R.B.G. thanks ...” Put sponsor acknowledgments in the unnumbered footnote on the first page.